

# Chris Addy

484-464-1249 | [chris.william.addy@gmail.com](mailto:chris.william.addy@gmail.com) | [github.com/chrisaddy](https://github.com/chrisaddy) | [linkedin.com/in/christopheraddy](https://linkedin.com/in/christopheraddy) | [aither.computer](https://aither.computer)

## PROJECTS

---

- Founder, aither.computer ([aither.computer](https://aither.computer))** 2024 — Present
- Building an ML observability and serving platform that unifies logging, deployment, and model explanation in a single developer-friendly interface.
  - Rust backend with Python SDK — combining high-performance async infrastructure with ergonomic APIs for ML practitioners.
  - Solving ML tooling fragmentation: one platform replacing the need for separate MLflow, model serving, and observability tools.

## WORK EXPERIENCE

---

**Engineering Manager, AI/ML** June 2025 — Present  
MNTN *Remote*

- Lead a team of 3 engineers while remaining hands-on, shipping 2 AI/ML products to production.
- Drove deliverability forecasting system to production, leading integration efforts for a model that predicts campaign budget spend and recommends targeting adjustments.
- Built media plan budget allocation engine from scratch, combining ML optimization with LLM-powered natural language explanations to help customers understand recommendations.
- Developed internal Vertex AI tooling and AI efficiency utilities that increased team shipping velocity from weekly to multiple PRs per engineer per day.

**Machine Learning Architect** January 2025 — June 2025  
URBN *Philadelphia, PA*

- Architected migration from Airflow to Flyte, improving developer experience with type-safe workflows and reducing pipeline debugging time for ML engineers.
- Redesigned MLflow architecture to standardize experiment tracking, enabling reproducible ML workflows and self-service model deployment for data scientists.
- Built internal ML platform tooling with Pandera and Pydantic for schema validation, giving engineers confidence in data pipeline correctness.
- Created DuckDB-based data processing utilities that simplified analyst workflows and eliminated pandas memory issues in production pipelines.

**Head of Machine Learning Engineering, Generative AI, eCommerce** May 2023 — December 2024  
Pepsico *Philadelphia, PA*

- Built ML platform that reduced pipeline development from months to days, enabling 3x more data scientists to ship models independently.
- Designed self-service ML infrastructure that expanded the number of analysts building production pipelines by 200%.
- Founded Generative AI Lab with 8 engineers, establishing internal tooling patterns for LLM integration and prompt engineering workflows.
- Re-architected MLOps lifecycle with standardized CI/CD templates and deployment automation, cutting model deployment time by 300%.
- Created internal Python SDK for ML workflows that became the standard interface for model training, evaluation, and serving.

**Principal Data Scientist & Machine Learning Engineer, eCommerce** July 2021 — May 2023  
Pepsico *Philadelphia, PA*

- Architected KubeFlow-based ML platform that standardized the path from data ingestion to model serving, reducing time-to-delivery by 50%.

- Built reusable pipeline components and templates that enabled rapid prototyping, cutting proof-of-concept time to three weeks.
- Designed unified data processing framework that standardized ingestion, transformation, and feature engineering across teams.
- Optimized AWS Batch infrastructure for ML training workloads, reducing compute costs and improving job scheduling efficiency.

**Senior & Lead Data Scientist**

Penn Interactive

June 2020 — April 2021

Philadelphia, PA

- Founded ML engineering function, building inference infrastructure and recommendation APIs that supported 75k users at launch.
- Developed ML serving infrastructure with RESTful APIs and real-time monitoring via DataDog dashboards.

**Data Scientist**

Heavywater

November 2018 — October 2019

Philadelphia, PA

- Built document classification pipeline handling 300+ document types with 96% accuracy, increasing processing throughput 10x.

---

**EDUCATION****Temple University***Master of Science, Statistics*

Philadelphia, PA

**Lehigh University***Bachelor of Arts, Economics*

Bethlehem, PA